

U4_AT2_Article Review Assessment

Decision Trees: From Efficient Prediction to Responsible AI

Paper link : <https://doi.org/10.3389/frai.2023.1124553>

Name of Student : B.BHARGAVI

Register Number: 192311136

Course & Code: CSA1711 Artificial Intelligence

Task 1: Article Summary

(a) Full Citation

Blockeel, H., Devos, L., Frénay, B., Nanfack, G., & Nijssen, S. (2023). Decision trees: from efficient prediction to responsible AI. *Frontiers in Artificial Intelligence*, 6, 1124553.

DOI: 10.3389/frai.2023.1124553

Domain / Application Area: Machine learning methodology — decision tree and tree-ensemble research, spanning predictive modeling, combinatorial optimization, and Responsible AI (fairness, robustness, explainability).

Specific ML Problem Addressed: The article does not target a single applied problem. It surveys how decision trees are learned, extended beyond classification and regression, optimized, constrained, and used as tools for building trustworthy AI systems.

(b) Key Objective

The authors aim to give a broad, chronological overview of decision tree research spanning roughly four decades, arguing that existing surveys are too narrow — each typically covers one specific use case (for example, only predictive accuracy or only one algorithmic family) rather than the full breadth of what decision trees

can do. The research gap they identify is that no single survey connects the classical view of decision trees as fast, interpretable predictors with their more recent and growing role in Responsible AI — ensuring models are fair, robust to adversarial attacks, and explainable. To fill this gap, the authors trace the field's evolution: starting with recursive partitioning and ensemble methods (bagging, boosting), moving through generalizations (multi-target, clustering, ranking, probabilistic trees), then combinatorial and gradient-based optimal tree learning, constraint integration, and finally dedicating a full section to fairness, adversarial robustness, and explainability. Their goal is to show that decision trees remain relevant not just for accuracy, but because their discrete, interpretable structure uniquely suits the demands of trustworthy, verifiable AI.

Task 2: Methodology Analysis

(a) ML Technique(s) Used

As this is a survey rather than a single-method paper, it covers a family of techniques rather than proposing or testing one:

- Core technique: Recursive partitioning (top-down induction) — the classic greedy algorithm underlying ID3, CART, and C4.5.
- Ensemble variants: Bagging, Random Forests, Gradient Boosting (XGBoost, LightGBM).
- Optimal tree learning: Exhaustive/combinatorial search (DL8, DL8.5, GOSDT), MILP-based methods (Bertsimas & Dunn), SAT/MaxSAT/Constraint Programming solvers.
- Gradient-based tree learning: Soft decision trees and oblique trees via differentiable objectives (e.g., neural decision forests).
- Specialized variants: Survival trees, fuzzy decision trees, Hoeffding trees (incremental/streaming), Isolation Forests (anomaly detection).

- Responsible AI extensions: Fairness-aware splitting heuristics, MILP-based robust/adversarial training (e.g., GROOT), and knowledge-distillation methods (LIME, TREPAN) for explainability.

Dataset Used — Size, Features, Source, Preparation: There is no single dataset, as this is a conceptual/methodological review rather than an empirical study. The paper instead references benchmark results and datasets used by the individual papers it surveys (for example, tabular benchmark datasets referenced via the Grinsztajn et al. 2022 comparison of tree-based models versus deep learning). No original data collection or preprocessing is performed by the authors.

(b) Mapping to CO4 Topics

CO4 Topic	Where It Maps in the Paper
Inductive Learning	Recursive partitioning as the core inductive bias — building general rules from labeled training pairs (Section 2.1–2.2)
Decision Trees	Entire Sections 2–5: variants, ensembles, optimal-tree search, constraint integration
Statistical Learning	Information gain/entropy-based splitting, variance-reduction heuristics, probability estimation trees, density estimation trees (Section 3)
Reinforcement Learning	Mentioned only briefly — decision trees used for RL policy verification/extraction

Methodological Strength: The review's structure is itself a strength — it is genuinely comprehensive, tracing decision trees from basic heuristics to cutting-edge optimal-tree solvers and Responsible AI applications, with a clear chronological and thematic scaffold that makes the field's evolution easy to follow.

Methodological Limitation: Because it is a “birds-eye view” rather than a systematic survey, the authors selectively cite illustrative papers rather than exhaustively covering the literature (a limitation they acknowledge). Coverage is therefore uneven — some sub-areas (e.g., evolutionary algorithms, Section 4.5) receive much thinner treatment than others (e.g., combinatorial optimization,

Section 4.3).

Task 3: Results & Critical Evaluation

(a) Key Results

No experiments are run by the authors themselves. What the paper reports are results summarised from other papers, primarily through two comparison tables:

- Table 1 in the article compares tree-ensemble verification methods (adversarial example generation and robustness checking) across exactness, anytime behavior, ability to generate concrete examples, code availability, and supported norms (l_0 , l_∞ , l_p).
- Table 2 in the article compares robust decision tree learning methods (e.g., RobustTrees, GROOT, TREANT, ROCT) across ensemble type, computational complexity, attack norm, and whether robustness guarantees are provided.

No accuracy, F1-score, or AUC-ROC values are reported directly — the paper remains at the level of qualitative and structural comparison (what a method guarantees, rather than what score it achieves).

Comparative Table (recreated from Table 2 of the article)

Method	Ensemble Type	Complexity	Norm	Robustness Guaranteed?
Adversarial boosting	GB	$n \log n$	l_0	No
RobustTrees	RF + GB	$n \log n$	l_∞	No
GROOT	RF	$n \log n$	l_∞	No
TREANT	RF	n^2	l_p	Yes
ROCT	Single tree	$\exp(n)$	l_∞	Yes
FPRDT	Single tree	$n \log n$	l_∞	Yes

(b) Critical Evaluation of Results

Since no primary experiments are conducted, the “realism” of metrics is not directly assessable — the paper's value lies in synthesis rather than in generating new empirical evidence.

Potential Bias: The authors are all decision-tree researchers (KU Leuven, Université de Namur, UCLouvain), so there is a natural selection bias toward papers that emphasise tree-based methods' strengths, with comparatively less critical scrutiny of where trees lose to deep learning (for example, on unstructured or high-dimensional data — acknowledged only briefly in Section 2.5).

What Would Strengthen It: A systematic (rather than illustrative) literature review methodology — such as PRISMA-style search and selection criteria — and inclusion of quantitative meta-analysis across the benchmark accuracy numbers reported by the individual cited papers, rather than only structural/qualitative comparison tables.

(c) Real-World Applicability

Deployment Areas: Tabular data problems such as finance, healthcare risk scoring, credit approval, and fraud detection — tree ensembles like XGBoost remain the practical default for structured data.

Responsible AI Relevance: Verified/robust trees are directly applicable to safety-critical domains (healthcare, air traffic control) where formal guarantees matter more than marginal accuracy gains.

Scaling Challenges:

- Data availability: Optimal-tree and MILP-based methods scale poorly with dataset size compared to heuristic recursive partitioning.
- Compute: Exact verification and optimal tree search are NP-hard, making them practical only for moderate tree/ensemble sizes.

- Ethical issues: Fairness-constrained tree learning still involves trade-offs between accuracy and fairness metrics that have no universally “correct” balance — this is ultimately a deployment/governance decision, not a purely technical one.

Task 4: Reflection & Learning Outcome

(a) Three Key Insights

Insight 1: Decision trees are inductive learners at their core, but the “informativeness” heuristic (entropy/information gain) that drives split choice is not guaranteed optimal — it is a greedy approximation. This connects to CO4's Inductive Learning topic, specifically the bias-variance trade-offs inherent in rule induction.

Insight 2: Optimal decision tree learning (DL8.5, GOSDT, MILP-based methods) reframes tree induction as a combinatorial/constraint-satisfaction problem rather than a purely statistical/heuristic one. This connects to CO4's Decision Trees topic, specifically the shift from heuristic to exact search-based learning.

Insight 3: Interpretability is not a single property but has at least four distinct forms — understanding the whole model, understanding feature importance, understanding a single prediction, and understanding the reasoning path. This connects to CO4's Statistical Learning / Explainability topic, adding nuance often missing from introductory coverage of “interpretable models.”

(b) Proposed Extension

Novel Experiment: Apply the fairness-aware optimal decision tree approaches discussed in Section 6.1 of the article (e.g., the DL8-style method of van der Linden et al., 2022) to a real-world Indian dataset — for example, a student performance or loan-approval dataset — to test whether the fairness-accuracy trade-off techniques generalise outside the demographic contexts (largely US/EU) typically used in the fairness-in-ML literature.

Feasibility: DL8.5 and similar solvers are open-source (as noted in the article) and computationally tractable for small-to-medium tabular datasets, making this well within scope for an academic project.

Expected Impact: This would test whether fairness metrics and constraint-aware heuristics (e.g., fair information gain) hold up across a different population distribution — directly relevant to responsible-AI deployment in the Indian context, an angle underexplored given the article's Western-dataset-heavy citation base.